


SWINBURNE

SWINBURNE UNIVERSITY OF TECHNOLOGY

# COS30019: Introduction to Artificial Intelligence



## Intelligent Agents



1

## Previously ...

- What is AI?
  - Four paradigms (think vs. act, human-like vs. rationally)
  - AI in movies/science fictions (food for thought)
  - AI in the real world
  - In COS30019, we will study “Systems that act rationally”
- “Systems that do the right thing”
  - The “Right thing” is the course of action that is expected to *maximize goal achievement given the available information.*




2

## Outline

- Agents and environments.
  - The vacuum-cleaner world
- The concept of rational behavior.
- Environments.
- Agent structure.

## Agents and environments

- Agents include human, robots, softbots, thermostats, etc.
- The *agent function* maps percept sequence to actions




- An agent can perceive its own actions, but not always its effects.

$$f : P^* \rightarrow A$$


## Agents and environments



- The *agent function* will internally be represented by the *agent program*.
- The agent program runs on the physical *architecture* to produce *f*.




## The vacuum-cleaner world – An example



- **Environment:** squares A and B
- **Percepts:** [location and content] e.g. [A, Dirty]
- **Actions:** left, right, suck, and no-op

## The vacuum-cleaner world – Agent function




Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
...	...



7

## The vacuum-cleaner world – An agent program



```

procedure REFLEX-VACUUM-AGENT ([location, status]) return an action
  if status == Dirty then return Suck
  else if location == A then return Right
  else if location == B then return Left
    
```

What is the right function? Can it be implemented in a small agent program?



8

## The concept of rationality



- A **rational agent** is one that does the right thing.
  - Every entry in the table is filled out correctly.
- What is the right thing?
  - Approximation: the most *successful* agent.
  - *Measure of success?*
- Performance measure should be **objective**
  - E.g. the amount of dirt cleaned within a certain time.
  - E.g. how clean the floor is.
  - ...
- *Performance measure according to what is wanted in the environment instead of how the agents should behave.*



9

## Rationality



- What is rational at a given time depends on four things:
  - Performance measure,
  - Prior environment knowledge,
  - Actions,
  - Percept sequence to date (sensors).
- DEF: *A rational agent chooses whichever action that maximizes the expected value of the performance measure given the percept sequence to date and prior environment knowledge.*



10

## Rationality



- Rationality  $\neq$  omniscience
  - An omniscient agent knows the actual outcome of its actions.
- Rationality  $\neq$  perfection
  - Rationality maximizes *expected* performance, while perfection maximizes *actual* performance.



11

## Rationality



- The proposed definition requires:
  - Information gathering/exploration
    - To maximize future rewards
  - Learn from percepts
    - Extending prior knowledge
  - Agent autonomy
    - Compensate for incorrect prior knowledge



12

## Is the vacuum cleaner agent rational?



- **Depend!**
- For example, **it's rational under the following assumptions:**
  - ☐ Performance measure: 1 point for each clean square over 'lifetime' of 1000 steps
  - ☐ 'geography' known but dirt distribution, initial position of agent not known
  - ☐ Clean squares stay clean, sucking cleans squares
  - ☐ Left and Right don't take agent outside environment
  - ☐ Available actions: Left, Right, Suck, NoOp
  - ☐ Agent knows where it is and whether that location contains dirt



13

## Environments



- To design a rational agent we must specify its **task environment**
- **PEAS** description of the task environment:
  - ☐ **P**erformance
  - ☐ **E**nvironment
  - ☐ **A**ctuators
  - ☐ **S**ensors



14

## Environments



### ■ E.g. Fully automated taxi:

#### ☐ PEAS description of the environment:

##### ☐ Performance

☐ Safety, destination, profits, legality, comfort

##### ☐ Environment

☐ Streets/freeways, other traffic, pedestrians, weather, ...

##### ☐ Actuators

☐ Steering, accelerating, brake, horn, speaker/display, ...

##### ☐ Sensors

☐ Video, sonar, speedometer, engine sensors, keyboard, GPS, ...



15

## Environment types



	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??				
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



16



## The game of backgammon



17

## Environment types



**Fully vs. partially observable:** an environment is full observable when the sensors can detect all aspects that are *relevant* to the choice of action.

	Crossword	Backgammon	Chess w/ clock	Taxi
<b>Observable??</b>				
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



18

## Environment types



**Fully vs. partially observable:** an environment is full observable when the sensors can detect all aspects that are *relevant* to the choice of action.


	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				

fully observable



partially observable

19



S, E

## Environment types



**Deterministic vs. stochastic:** if the next environment state is completely determined by the current state the executed action then the environment is deterministic.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



20

## Environment types



**Deterministic vs. stochastic:** if the next environment state is completely determined by the current state the executed action then the environment is deterministic.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??				
Static??				
Discrete??				
Single-agent??				

deterministic




stochastic

21



actuator



## Environment types



**Episodic vs. sequential:** In an episodic environment the agent's experience can be divided into atomic steps where the agents perceives and then performs A single action. The choice of action depends only on the episode itself

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??				
Static??				
Discrete??				
Single-agent??				



22

100%  
A, E

60%

## Environment types



**Episodic vs. sequential:** In an episodic environment the agent's experience can be divided into atomic steps where the agent perceives and then performs a single action. The choice of action depends only on the episode itself

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??				
Discrete??				
Single-agent??				

episodic

sequential



23

actuator

affects

## Environment types



**Static vs. dynamic:** If the environment can change while the agent is choosing an action, the environment is dynamic. Semi-dynamic if the agent's performance changes even when the environment remains the same.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??				
Discrete??				
Single-agent??				



24

## Environment types




**Static vs. dynamic:** If the environment can change while the agent is choosing an action, the environment is dynamic. Semi-dynamic if the agent's performance changes even when the environment remains the same.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??				
Single-agent??				

static

semi-dynamic

dynamic



## Environment types



**Discrete vs. continuous:** This distinction can be applied to the state of the environment, the way time is handled and to the percepts/actions of the agent.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??				
Single-agent??				



26

## Environment types




**Discrete vs. continuous:** This distinction can be applied to the state of the environment, the way time is handled and to the percepts/actions of the agent.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??				

discrete

continuous



## Environment types



**Single vs. multi-agent:** Does the environment contain other agents who are also maximizing some performance measure that depends on the current agent's actions?

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??				



28

## Environment types



**Single vs. multi-agent:** Does the environment contain other agents who are also maximizing some performance measure that depends on the current agent's actions?

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??	YES	NO	NO	NO

multi-agent

single-agent



## Environment types



- The simplest environment is
  - ☐ Fully observable, deterministic, episodic, static, discrete and single-agent.
- Most real situations are:
  - ☐ Partially observable, stochastic, sequential, dynamic, continuous and multi-agent.



30

## Agent types



- How does the inside of the agent work?
  - Agent = architecture + program
- All agents have the same skeleton:
  - Input = current percepts
  - Output = action
  - Program = manipulates input to produce output
- Note difference with agent function.



31

## Agent types



Function `TABLE-DRIVEN_AGENT(percept)` **returns** an action

**static:** *percepts*, a sequence initially empty

*table*, a table of actions, indexed by percept sequence

append *percept* to the end of *percepts*

*action* ← LOOKUP(*percepts*, *table*)

**return** *action*

**This approach is doomed to failure**



32



## Agent types

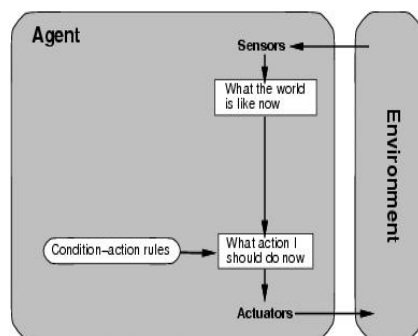


- Four basic kind of agent programs will be discussed:
  - ☐ Simple reflex agents
  - ☐ Model-based reflex agents
  - ☐ Goal-based agents
  - ☐ Utility-based agents
- All these can be turned into learning agents.
  - ☐ And that gives you four additional advanced agent types



33

## Agent types; simple reflex




- Select action on the basis of *only the current* percept.
  - ☐ E.g. the vacuum-agent
- Large reduction in possible percept/action situations(next page).
- Implemented through *condition-action rules*
  - ☐ If dirty then suck



34

## The vacuum-cleaner world



```
function REFLEX-VACUUM-AGENT ([location, status]) return an action
  if status == Dirty then return Suck
  else if location == A then return Right
  else if location == B then return Left
```

Reduction from  $4^T$  to 4 entries



35

## Agent types; simple reflex



```
function SIMPLE-REFLEX-AGENT(percept) returns an action
```

**static:** *rules*, a set of condition-action rules

```
state ← INTERPRET-INPUT(percept)
```

```
rule ← RULE-MATCH(state, rules)
```

```
action ← RULE-ACTION[rule]
```

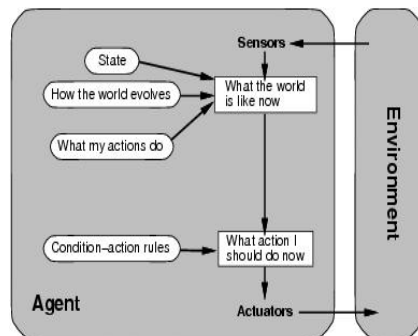
```
return action
```

Will only work if the environment is *fully observable*  
otherwise infinite loops may occur.



36

## Agent types; reflex and state



- To tackle *partially observable* environments.

- Maintain internal state

- Over time update state using world knowledge

- How does the world change.

- How do actions affect world.

⇒ *Model of World*



37

## Agent types; reflex and state



**function** REFLEX-AGENT-WITH-STATE(*percept*) **returns** an action

**static:** *rules*, a set of condition-action rules

*state*, a description of the current world state

*actions*, the most recent actions.

*state* ← UPDATE-STATE(*state*, *actions*, *percept*)

*rule* ← RULE-MATCH(*state*, *rules*)

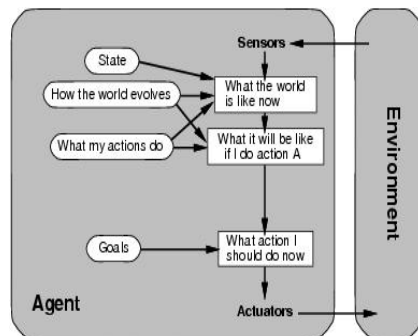
*action* ← RULE-ACTION[*rule*]

**return** *action*



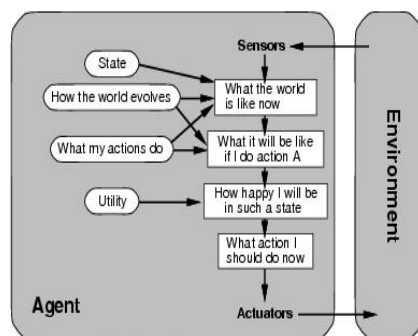
38

## Agent types; goal-based



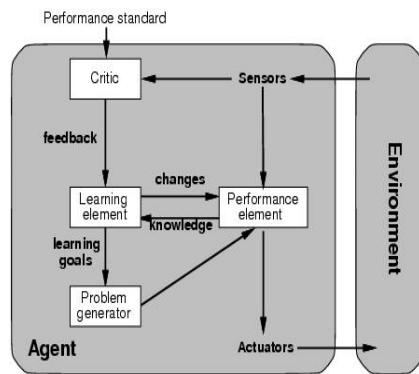
- The agent needs a goal to know which situations are *desirable*.
  - Things become difficult when long sequences of actions are required to find the goal.
- Typically investigated in **search** and **planning** research.
- Major difference: future is taken into account
- Is more flexible since knowledge is represented explicitly and can be manipulated.

## Agent types; utility-based



- Certain goals can be reached in different ways.
  - Some are better, have a higher utility.
- Utility function maps a (sequence of) state(s) onto a real number.
- Improves on goals:
  - Selecting between conflicting goals
  - Select appropriately between several goals based on likelihood of success.

## Agent types; learning

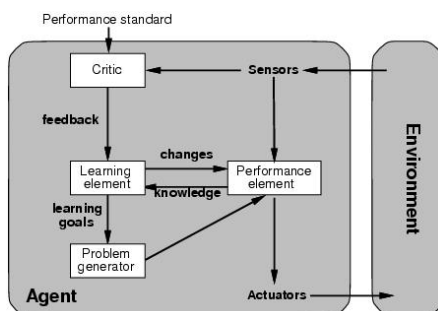


41

- All previous agent-programs describe methods for selecting *actions*.

- ☐ Yet it does not explain the origin of these programs.
- ☐ Learning mechanisms can be used to perform this task.
- ☐ Teach them instead of instructing them.
- ☐ Advantage is the robustness of the program toward initially unknown environments.

## Agent types; learning



42

- *Learning element*: introduce improvements in performance element.
  - ☐ Critic provides feedback on agents performance based on fixed performance standard.
- *Performance element*: selecting actions based on percepts.
  - ☐ Corresponds to the previous agent programs
- *Problem generator*: suggests actions that will lead to new and informative experiences.
  - ☐ Exploration vs. exploitation

## Summary: Agents



- An **agent** perceives and acts in an environment, has an architecture, and is implemented by an agent program.
- Task environment – **PEAS (Performance, Environment, Actuators, Sensors)**
- An **ideal agent** always chooses the action which maximizes its expected performance, given its percept sequence so far.
- An **autonomous learning agent** uses its own experience rather than built-in knowledge of the environment by the designer.
- An **agent program** maps from percept to action and updates internal state.
  - ☐ **Reflex agents** respond immediately to percepts.
  - ☐ **Goal-based agents** act in order to achieve their goal(s).
  - ☐ **Utility-based agents** maximize their own utility function.
- **Representing knowledge** is important for successful agent design.
- The most challenging environments are not fully observable, nondeterministic, dynamic, and continuous

